

Approfondimento di Intelligenza Artificiale


Tristano Munini

11 ottobre 2020


Introduzione

TODO


template



template



Reparametrizing the sampling layer



Key Idea:

$$z \sim \mathcal{N}(\mu, \sigma^2)$$

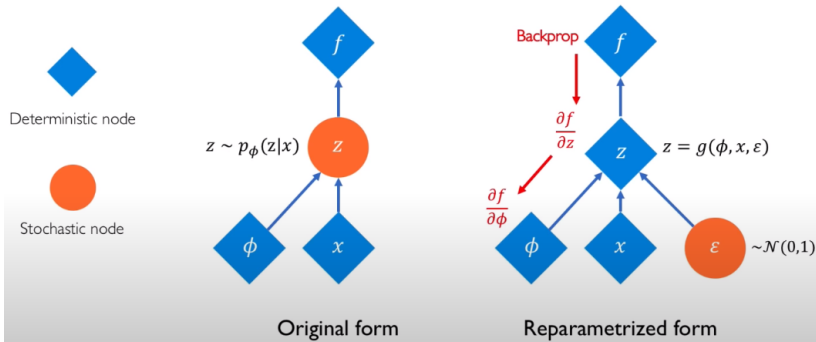
Consider the sampled latent vector z as a sum of

- a fixed μ vector,
- and fixed σ vector, scaled by random constants drawn from the prior distribution

$$\Rightarrow z = \mu + \sigma \odot \epsilon$$

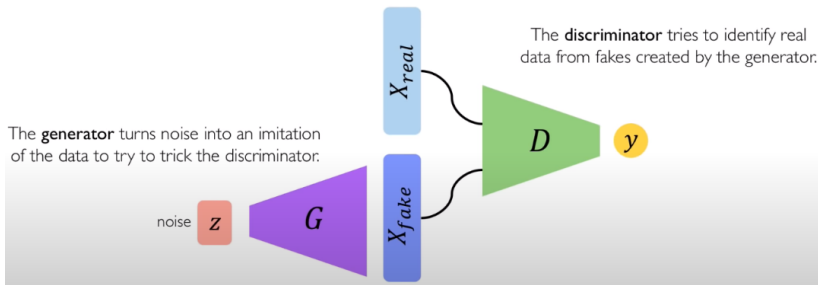
where $\epsilon \sim \mathcal{N}(0,1)$

Reparametrizing the sampling layer




Generative Adversarial Networks (GANs)

Generative Adversarial Networks (GANs) are a way to make a generative model by having two neural networks compete with each other:



template



Discounted Total Reward (Return)

$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i = \gamma^t r_t + \gamma^{t+1} r_{t+1} \dots + \gamma^{t+n} r_{t+n} + \dots$$

γ : discount factor; $0 < \gamma < 1$

template

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

Total reward, R_t , is the discounted sum of all rewards obtained from time t

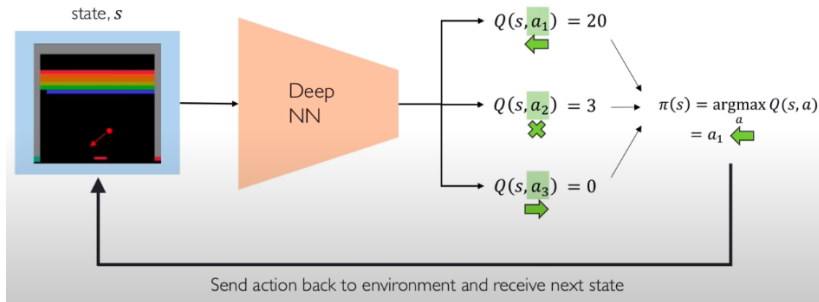
$$Q(s_t, a_t) = \mathbb{E}[R_t | s_t, a_t]$$

The Q-function captures the **expected total future reward** an agent in **state, s** , can receive by executing a certain **action, a**


$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

template



Use NN to learn Q-function and then use to infer the optimal policy, $\pi(s)$



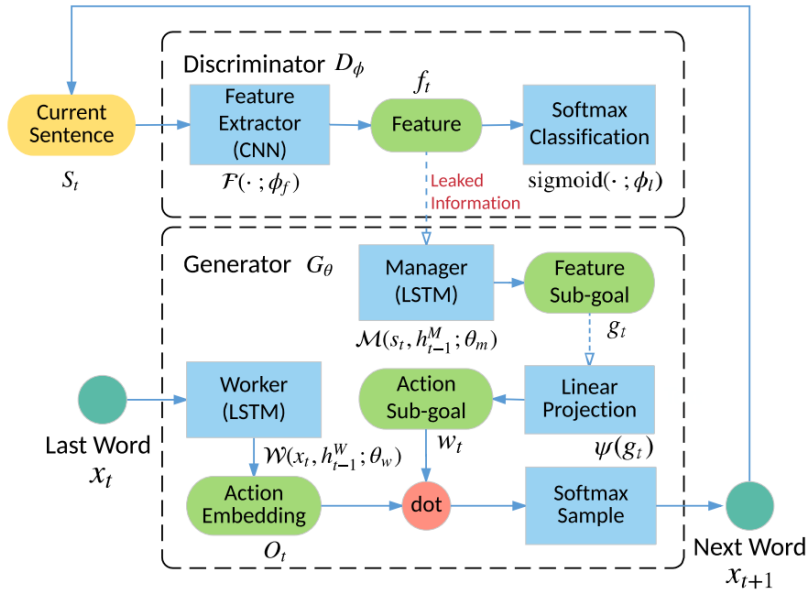
template



template



template



template

