

Approfondimento di Intelligenza Artificiale

Tristano Munini

11 ottobre 2020

Indice


- ▶ AutoEncoders (AE) e Variational AutoEncoders (VAE)
- ▶ Generative Adversarial Network (GAN)
- ▶ Reinforcement Learning (RL)
 - ▶ QNN
 - ▶ Policy Gradient
- ▶ GAN + RL per generazione di testi
 - ▶ SeqGAN
 - ▶ LeakGAN

AutoEncoders

$x \in \text{real data space}$

$z \in \text{latent space}$

$\hat{x} \in \text{real data space}$

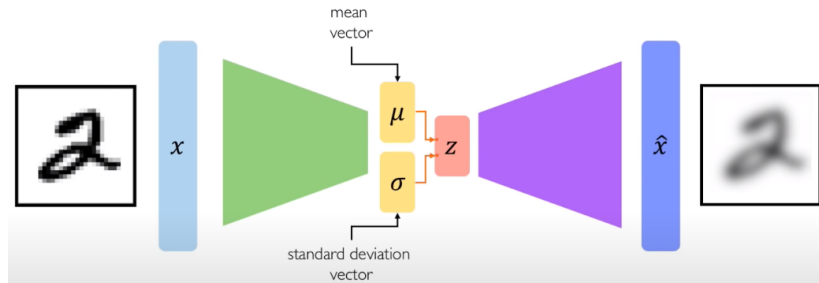


The point of maximum compression, where z is produced, is called bottleneck of the network

Variational AutoEncoders

z is sampled from the distribution with mean vector μ and standard deviation vector σ .

We want an \hat{x} that is similar to x but that has some differences.
We want a variation of x .



Reparametrization Trick

We can not use $z \sim \mathcal{N}(\mu, \sigma^2)$

So we use

$$z = \mu + \sigma \odot \varepsilon \quad \text{where } \varepsilon \sim \mathcal{N}(0, 1)$$

Where \odot notes the element-wise multiplication

Now z is the sum of two fixed vectors, μ and σ , and a random constant ε used as a weight

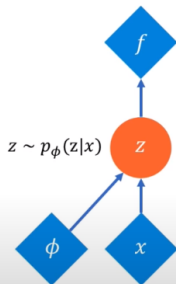
Reparametrization Trick




Deterministic node



Stochastic node



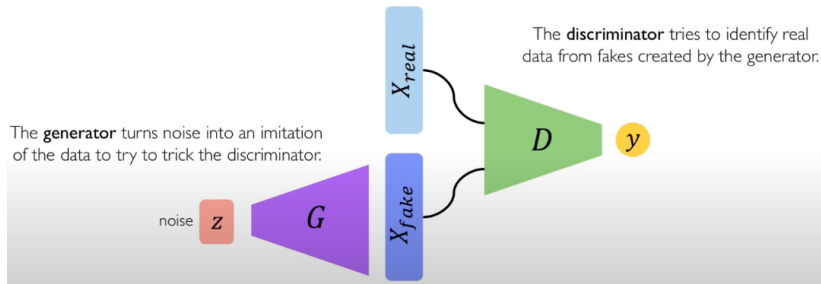
Original form



Reparametrized form

Generative Adversarial Networks

GANs are a way to make a generative model by having two neural networks compete with each other



z can be sampled from $\mathcal{N}(0, 1)$ as in the VAEs

GAN Training

It is a min-max game


$$\begin{aligned} \min_G \max_D V(D, G) = & \mathbb{E}_{x \sim P_{real(x)}} [\log D(x)] \\ & + \mathbb{E}_{z \sim P_z(x)} [\log(1 - D(G(z)))] \end{aligned}$$

Where $P_{real(x)}$ is the probability distribution of the real data and $P_z(x) = \mathcal{N}(0, 1)$ in our case

GAN Problems

- ▶ Hard to converge to a good solution
- ▶ Vanishing gradient problem
- ▶ Model collapse
- ▶ Hard to find equilibrium

template



Discounted Total Reward (Return)

$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i = \gamma^t r_t + \gamma^{t+1} r_{t+1} \dots + \gamma^{t+n} r_{t+n} + \dots$$

γ : discount factor; $0 < \gamma < 1$

template

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

Total reward, R_t , is the discounted sum of all rewards obtained from time t


$$Q(s_t, a_t) = \mathbb{E}[R_t | s_t, a_t]$$

The Q-function captures the **expected total future reward** an agent in **state, s** , can receive by executing a certain **action, a**


$$\pi^*(s) = \operatorname{argmax}_a Q(s, a)$$

template



Use NN to learn Q-function and then use to infer the optimal policy, $\pi(s)$




template



template



template



template

